



Phase 1 — Setup (one-time per dataset)

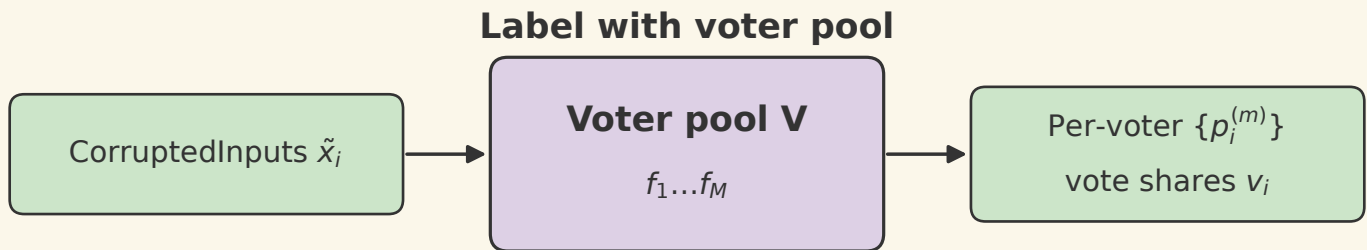


for each (corruption, severity)

Phase 2 — Generate one benchmark setting B_τ



each image logged with RNG seed, parameters, and output sha256



Release benchmark B_τ

$$B_\tau = \{x_i, \tilde{x}_i, y_i, \tilde{y}_i, v_i, \{p_i^{(m)}\}, \tau = (c, \ell)\}$$

clean input · corrupted input · ground-truth label · sampled training label · vote shares · per-voter dists · corruption metadata

optional CILN-C filter: keep only images all voters classify correctly on the clean input